

```
import pandas as pd
import os
from google.colab import drive, files
drive.mount('/content/drive')
uploaded = files.upload()
```

```
files = os.listdir()
print(files)
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

<IPython.core.display.HTML object>

```
Saving diabetes.csv to diabetes (1).csv
['.config', 'diabetes.csv', 'diabetes (1).csv', 'drive',
'sample_data']
```

```
!pip install seaborn
```

```
Requirement already satisfied: seaborn in
/usr/local/lib/python3.10/dist-packages (0.12.2)
Requirement already satisfied: numpy!=1.24.0,>=1.17 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (1.23.5)
Requirement already satisfied: pandas>=0.25 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (1.5.3)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (1.2.0)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (4.46.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (1.4.5)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (23.2)
Requirement already satisfied: pillow>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (3.1.1)
Requirement already satisfied: python-dateutil>=2.7 in
```

```

/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.1-
>seaborn) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=0.25->seaborn)
(2023.3.post1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7-
>matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)

```

```
dia = 'diabetes.csv'
```

```
diabetes = pd.read_csv(dia, sep = ",")
print(diabetes)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
[768 rows x 9 columns]
```

```
diabetes.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
diabetes.columns.values
```

```
array(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age',  
      'Outcome'],  
      dtype=object)
```

```
diabetes.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000

max	17.000000	199.000000	122.000000	99.000000
846.000000				

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
diabetes.nunique()
```

Pregnancies	17
Glucose	136
BloodPressure	47
SkinThickness	51
Insulin	186
BMI	248
DiabetesPedigreeFunction	517
Age	52
Outcome	2

dtype: int64

Monotonih atributa ovdje nemamo, ali najbliže monotonom atributu je DiabetesPedigreeFunction koji ima 517 različitih vrijednosti.

```
diabetes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                             768 non-null    int64
2   BloodPressure                       768 non-null    int64
3   SkinThickness                      768 non-null    int64
4   Insulin                            768 non-null    int64
5   BMI                                768 non-null    float64
6   DiabetesPedigreeFunction            768 non-null    float64
7   Age                                768 non-null    int64
8   Outcome                            768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

diabetes.isna().sum()
```

```

Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64

```

```
diabetes.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000
max	17.000000	199.000000	122.000000	99.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

Stršeće značajke bi bile značajke Pregnancies, SkinThickness i Insulin. Također je neobično to što značajke Glucose, BloodPressure, SkinThickness, Insulin i BMI imaju 0 kao minimalnu vrijednost.

```

Atributi = ["Glucose", "BloodPressure", "SkinThickness", "Insulin",
            "BMI"]
diabetes2 = diabetes[(diabetes[Atributi] != 0).all(axis=1)]
diabetes2.describe()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	392.000000	392.000000	392.000000	392.000000
mean	3.301020	122.627551	70.663265	29.145408
std	3.211424	30.860781	12.496092	10.516424
min	0.000000	56.000000	24.000000	7.000000
25%	1.000000	99.000000	62.000000	21.000000
50%	2.000000	119.000000	70.000000	29.000000
75%	5.000000	143.000000	78.000000	37.000000
max	17.000000	198.000000	110.000000	63.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	392.000000	392.000000	392.000000	392.000000
mean	33.086224	0.523046	30.864796	0.331633
std	7.027659	0.345488	10.200777	0.471401
min	18.200000	0.085000	21.000000	0.000000
25%	28.400000	0.269750	23.000000	0.000000
50%	33.200000	0.449500	27.000000	0.000000
75%	37.100000	0.687000	36.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
import matplotlib.pyplot as plt
fi, ax = plt.subplots(4, 2)
fi.set_figheight(17)
fi.set_figwidth(17)
fi.add_subplot(4, 2, 1)
plt.hist(diabetes.Pregnancies, bins = 10)
plt.xlabel('Pregnancies')
plt.ylabel('Count')
plt.axvline(diabetes.Pregnancies.mean(), color = 'red', label = 'mean')
plt.axvline(diabetes.Pregnancies.median(), color = 'g', label = 'median')
plt.hist(diabetes[diabetes.Outcome == 1].Pregnancies, bins = 10, range = (0, 17), label='Pregnancies w/ diabetes', color='red', alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].Pregnancies, bins = 10, range = (0, 17), label='Pregnancies w/o diabetes', color='green', alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 2)
plt.hist(diabetes.SkinThickness, bins = 10)
plt.xlabel('SkinThickness')
plt.ylabel('Count')
```

```

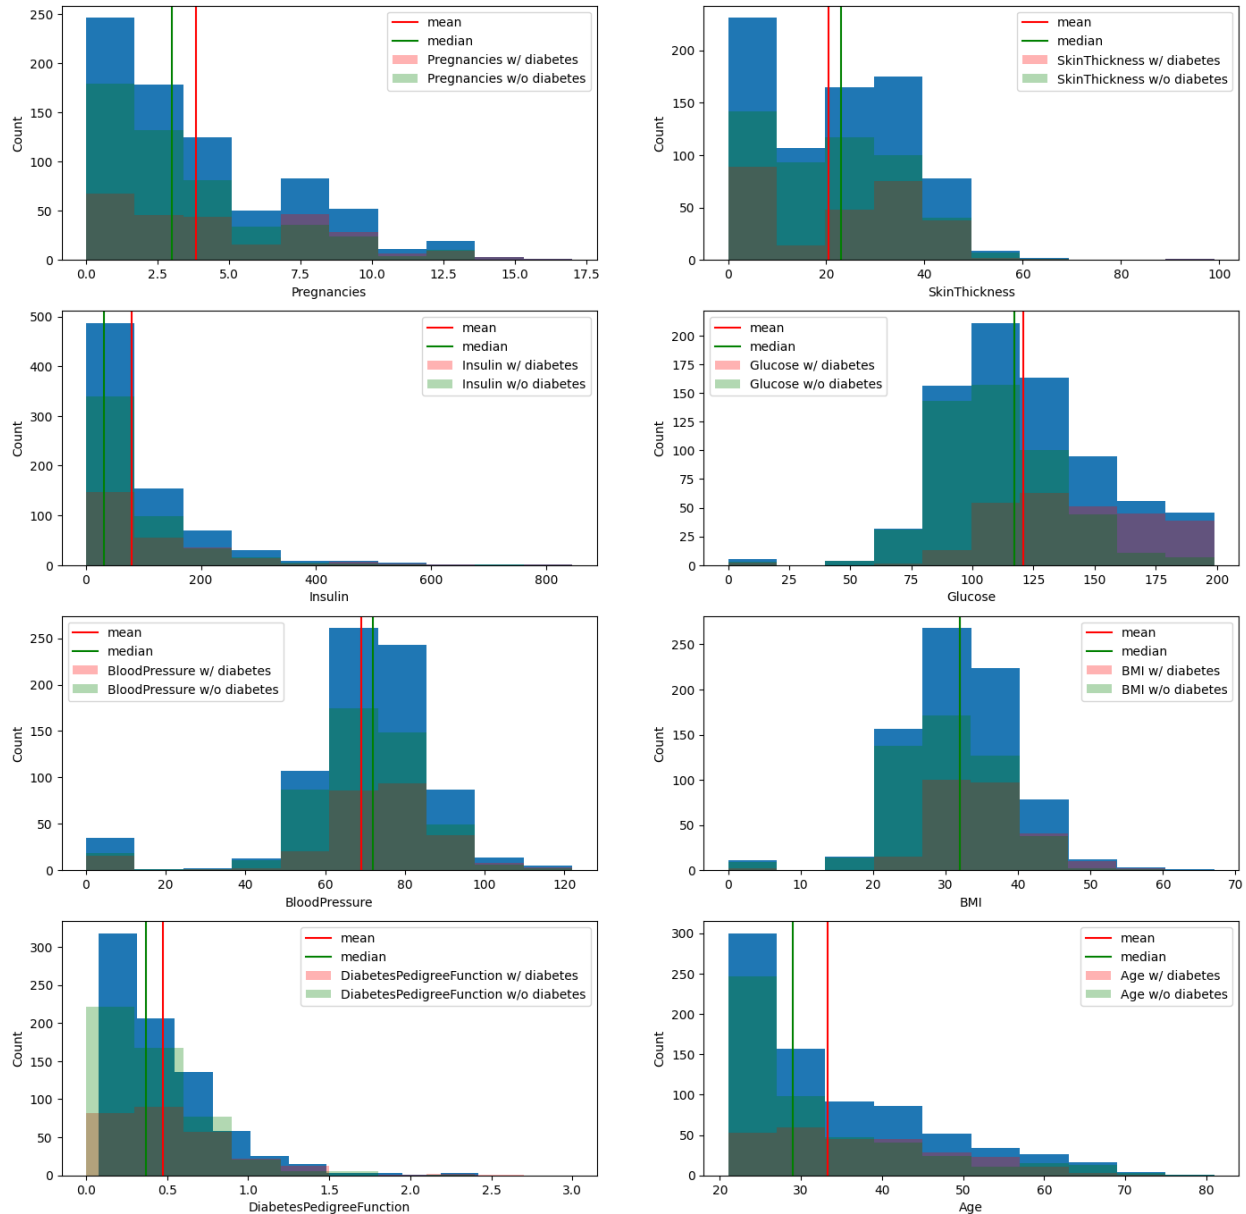
plt.axvline(diabetes.SkinThickness.mean(), color = 'red', label =
'mean')
plt.axvline(diabetes.SkinThickness.median(), color = 'g', label =
'median')
plt.hist(diabetes[diabetes.Outcome == 1].SkinThickness, bins = 10,
range = (0, 99), label='SkinThickness w/ diabetes', color='red',
alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].SkinThickness, bins = 10,
range = (0, 99), label='SkinThickness w/o diabetes', color='green',
alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 3)
plt.hist(diabetes.Insulin, bins = 10)
plt.xlabel('Insulin')
plt.ylabel('Count')
plt.axvline(diabetes.Insulin.mean(), color = 'red', label = 'mean')
plt.axvline(diabetes.Insulin.median(), color = 'g', label = 'median')
plt.hist(diabetes[diabetes.Outcome == 1].Insulin, bins = 10, range =
(0, 846), label='Insulin w/ diabetes', color='red', alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].Insulin, bins = 10, range =
(0, 846), label='Insulin w/o diabetes', color='green', alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 4)
plt.hist(diabetes.Glucose, bins = 10)
plt.xlabel('Glucose')
plt.ylabel('Count')
plt.axvline(diabetes.Glucose.mean(), color = 'red', label = 'mean')
plt.axvline(diabetes.Glucose.median(), color = 'g', label = 'median')
plt.hist(diabetes[diabetes.Outcome == 1].Glucose, bins = 10, range =
(0, 199), label='Glucose w/ diabetes', color='red', alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].Glucose, bins = 10, range =
(0, 199), label='Glucose w/o diabetes', color='green', alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 5)
plt.hist(diabetes.BloodPressure, bins = 10)
plt.xlabel('BloodPressure')
plt.ylabel('Count')
plt.axvline(diabetes.BloodPressure.mean(), color = 'red', label =
'mean')
plt.axvline(diabetes.BloodPressure.median(), color = 'g', label =
'median')
plt.hist(diabetes[diabetes.Outcome == 1].BloodPressure, bins = 10,
range = (0, 122), label='BloodPressure w/ diabetes', color='red',
alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].BloodPressure, bins = 10,
range = (0, 122), label='BloodPressure w/o diabetes', color='green',
alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 6)

```

```

plt.hist(diabetes.BMI, bins = 10)
plt.xlabel('BMI')
plt.ylabel('Count')
plt.axvline(diabetes.BMI.mean(), color = 'red', label = 'mean')
plt.axvline(diabetes.BMI.median(), color = 'g', label = 'median')
plt.hist(diabetes[diabetes.Outcome == 1].BMI, bins = 10, range = (0,
67), label='BMI w/ diabetes', color='red', alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].BMI, bins = 10, range = (0,
67), label='BMI w/o diabetes', color='green', alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 7)
plt.hist(diabetes.DiabetesPedigreeFunction, bins = 10)
plt.xlabel('DiabetesPedigreeFunction')
plt.ylabel('Count')
plt.axvline(diabetes.DiabetesPedigreeFunction.mean(), color = 'red',
label = 'mean')
plt.axvline(diabetes.DiabetesPedigreeFunction.median(), color = 'g',
label = 'median')
plt.hist(diabetes[diabetes.Outcome == 1].DiabetesPedigreeFunction,
bins = 10, range = (0, 3), label='DiabetesPedigreeFunction w/
diabetes', color='red', alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].DiabetesPedigreeFunction,
bins = 10, range = (0, 3), label='DiabetesPedigreeFunction w/o
diabetes', color='green', alpha=0.3)
plt.legend()
fi.add_subplot(4, 2, 8)
plt.hist(diabetes.Age, bins = 10)
plt.xlabel('Age')
plt.ylabel('Count')
plt.axvline(diabetes.Age.mean(), color = 'red', label = 'mean')
plt.axvline(diabetes.Age.median(), color = 'g', label = 'median')
plt.hist(diabetes[diabetes.Outcome == 1].Age, bins = 10, range = (21,
81), label='Age w/ diabetes', color='red', alpha=0.3)
plt.hist(diabetes[diabetes.Outcome == 0].Age, bins = 10, range = (21,
81), label='Age w/o diabetes', color='green', alpha=0.3)
plt.legend()
for a in ax.flatten():
    a.set_xticks([])
    a.set_yticks([])
    a.set_xticklabels([])
    a.set_yticklabels([])
    a.spines['top'].set_visible(False)
    a.spines['right'].set_visible(False)
    a.spines['bottom'].set_visible(False)
    a.spines['left'].set_visible(False)

```

```

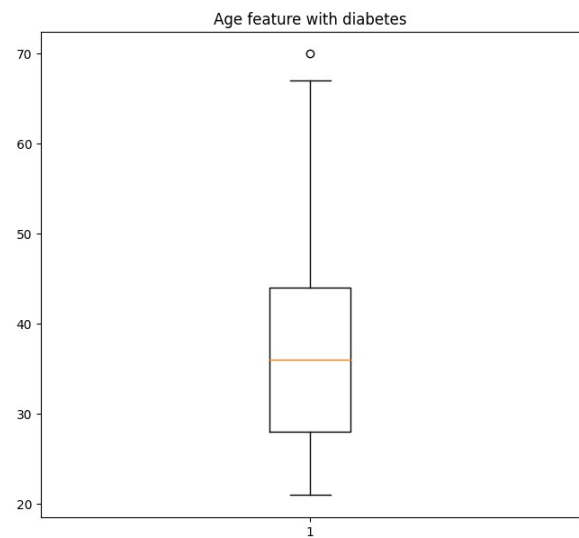
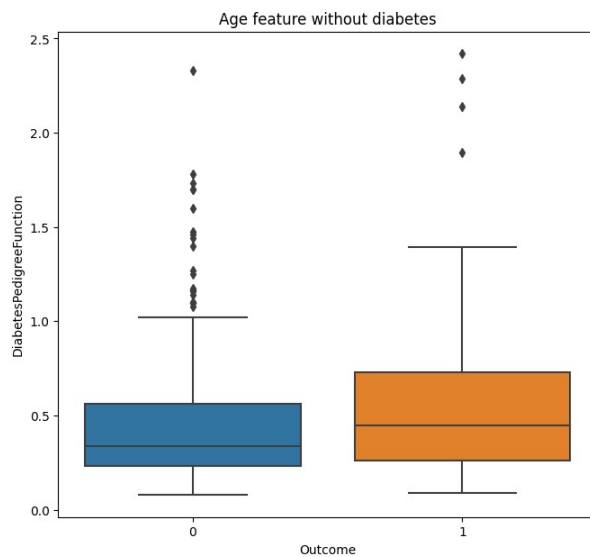
fi, ax = plt.subplots(1, 2)
fi.set_figheight(7)
fi.set_figwidth(17)
fi.add_subplot(1, 2, 1)
sns.boxplot(x = diabetes["Outcome"], y =
diabetes["DiabetesPedigreeFunction"])
plt.title("Age feature without diabetes")
fi.add_subplot(1, 2, 2)
plt.boxplot(diabetes[diabetes.Outcome == 1].Age)
plt.title("Age feature with diabetes")
for a in ax.flatten():
    a.set_xticks([])
    a.set_yticks([])

```

```

a.set_xticklabels([])
a.set_yticklabels([])
a.spines['top'].set_visible(False)
a.spines['right'].set_visible(False)
a.spines['bottom'].set_visible(False)
a.spines['left'].set_visible(False)
plt.show()

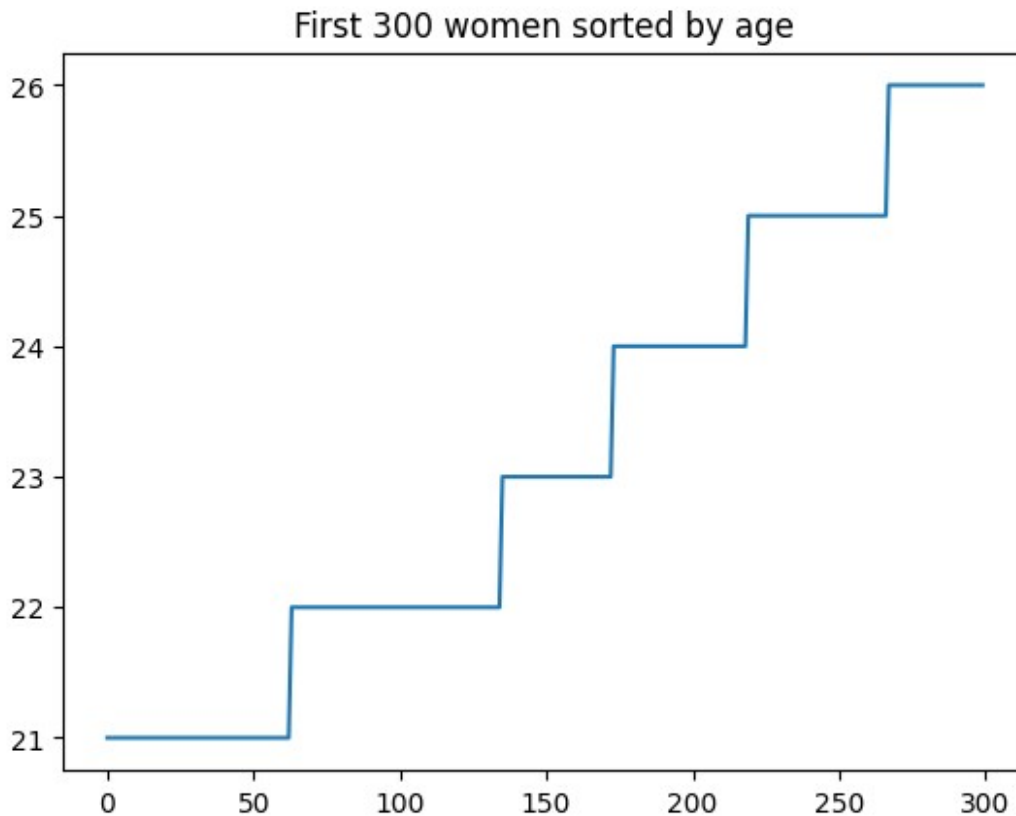
```







```

tmp = diabetes.copy()
sorted_age = tmp.sort_values(by=['Age']).Age.values[:300]
person = [x for x in range(300)]
plt.plot(person, sorted_age)
plt.title("First 300 women sorted by age")
plt.show()

```



```
fi, ax = plt.subplots(2, 2)
fi.set_figheight(17)
fi.set_figwidth(17)
fi.add_subplot(2, 2, 1)
plt.scatter(diabetes2.Age, diabetes2.Pregnancies)
fi.add_subplot(2, 2, 2)
plt.scatter(diabetes2.BMI, diabetes2.BloodPressure)
fi.add_subplot(2, 2, 3)
plt.scatter(diabetes2.Glucose, diabetes2.Insulin)
fi.add_subplot(2, 2, 4)
plt.scatter(diabetes2.DiabetesPedigreeFunction, diabetes2.Age)
for a in ax.flatten():
    a.set_xticks([])
    a.set_yticks([])
    a.set_xticklabels([])
    a.set_yticklabels([])
    a.spines['top'].set_visible(False)
    a.spines['right'].set_visible(False)
    a.spines['bottom'].set_visible(False)
    a.spines['left'].set_visible(False)
plt.show()
```



```





fi, ax = plt.subplots(2, 2)
fi.set_figheight(17)
fi.set_figwidth(17)
mask_survived = diabetes.Outcome == 1
fi.add_subplot(2, 2, 1)
plt.scatter(diabetes2.loc[mask_survived, 'Age'],
diabetes2.loc[mask_survived, 'BloodPressure'], alpha=0.5, c='r',
label='Diabetes')
plt.scatter(diabetes2.loc[~mask_survived, 'Age'],
diabetes2.loc[~mask_survived, 'BloodPressure'], alpha=0.5, c='g',
label='No diabetes')
plt.xlabel('Age')
plt.ylabel('Blood Pressure')

```


```

plt.grid()
plt.legend(loc='best')
fi.add_subplot(2, 2, 2)
plt.scatter(diabetes2.loc[mask_survived, 'SkinThickness'],
diabetes2.loc[mask_survived, 'Glucose'], alpha=0.5, c='r',
label='Diabetes')
plt.scatter(diabetes2.loc[~mask_survived, 'SkinThickness'],
diabetes2.loc[~mask_survived, 'Glucose'], alpha=0.5, c='g', label='No
diabetes')
plt.xlabel('SkinThickness')
plt.ylabel('Glucose')
plt.grid()
plt.legend(loc='best')
fi.add_subplot(2, 2, 3)
plt.scatter(diabetes2.loc[mask_survived, 'Insulin'],
diabetes2.loc[mask_survived, 'BMI'], alpha=0.5, c='r',
label='Diabetes')
plt.scatter(diabetes2.loc[~mask_survived, 'Insulin'],
diabetes2.loc[~mask_survived, 'BMI'], alpha=0.5, c='g', label='No
diabetes')
plt.xlabel('Insulin')
plt.ylabel('BMI')
plt.grid()
plt.legend(loc='best')
fi.add_subplot(2, 2, 4)
plt.scatter(diabetes2.loc[mask_survived, 'BloodPressure'],
diabetes2.loc[mask_survived, 'DiabetesPedigreeFunction'], alpha=0.5,
c='r', label='Diabetes')
plt.scatter(diabetes2.loc[~mask_survived, 'BloodPressure'],
diabetes2.loc[~mask_survived, 'DiabetesPedigreeFunction'], alpha=0.5,
c='g', label='No diabetes')
plt.xlabel('BloodPressure')
plt.ylabel('DiabetesPedigreeFunction')
plt.grid()
plt.legend(loc='best')
plt.show()

```



```
import seaborn as sns
diabetes_cor = diabetes2.loc[:,diabetes2.columns.isin(['Pregnancies',
'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
'DiabetesPedigreeFunction'])]
sns.heatmap(diabetes_cor.corr())
plt.show()
```



```
sns.set(rc={'figure.figsize':(20,10)})
sns.heatmap(diabetes_cor.corr(), annot=True, fmt=".2f", cmap="YlGnBu",
linewidths=0.5)
plt.show()
```



```
sns.set(rc={'figure.figsize':(25,15)})
sns.pairplot(diabetes2.loc[:, [ 'Pregnancies', 'Glucose',
'BloodPressure', 'SkinThickness', 'Insulin', 'BMI',
'DiabetesPedigreeFunction', 'Age', 'Outcome']], hue="Outcome")
plt.show()
```




```
diabetes.isnull().values.any()
```

```
False
```

```
Atributi = ["Glucose", "BloodPressure", "SkinThickness", "Insulin",
"BMI"]
sred = diabetes[Atributi].mean()
diabetes[Atributi] = diabetes[Atributi].mask(diabetes[Atributi] == 0,
pd.NA)
diabetes[Atributi] = diabetes[Atributi].apply(lambda col:
col.fillna(sred[col.name]))
diabetes.describe()
```

```
      Pregnancies      Glucose  BloodPressure  SkinThickness
Insulin \
```

count	768.000000	768.000000	768.000000	768.000000
768.000000				
mean	3.845052	121.681605	72.254807	26.606479
118.660163				
std	3.369578	30.436016	12.115932	9.631241
93.080358				
min	0.000000	44.000000	24.000000	7.000000
14.000000				
25%	1.000000	99.750000	64.000000	20.536458
79.799479				
50%	3.000000	117.000000	72.000000	23.000000
79.799479				
75%	6.000000	140.250000	80.000000	32.000000
127.250000				
max	17.000000	199.000000	122.000000	99.000000
846.000000				

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.450805	0.471876	33.240885	0.348958
std	6.875374	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
Atributi = ["Glucose", "BloodPressure", "SkinThickness", "Insulin",
" BMI"]
sred = diabetes[Atributi].median()
diabetes[Atributi] = diabetes[Atributi].mask(diabetes[Atributi] == 0,
pd.NA)
diabetes[Atributi] = diabetes[Atributi].apply(lambda col:
col.fillna(sred[col.name]))
diabetes.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
768.000000				
mean	3.845052	121.681605	72.254807	26.606479
118.660163				
std	3.369578	30.436016	12.115932	9.631241
93.080358				
min	0.000000	44.000000	24.000000	7.000000
14.000000				
25%	1.000000	99.750000	64.000000	20.536458
79.799479				
50%	3.000000	117.000000	72.000000	23.000000
79.799479				

75%	6.000000	140.250000	80.000000	32.000000
127.250000				
max	17.000000	199.000000	122.000000	99.000000
846.000000				

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.450805	0.471876	33.240885	0.348958
std	6.875374	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
import seaborn as sns
sns.pairplot(diabetes, hue="Outcome")
<seaborn.axisgrid.PairGrid at 0x7c95a0a69150>
```

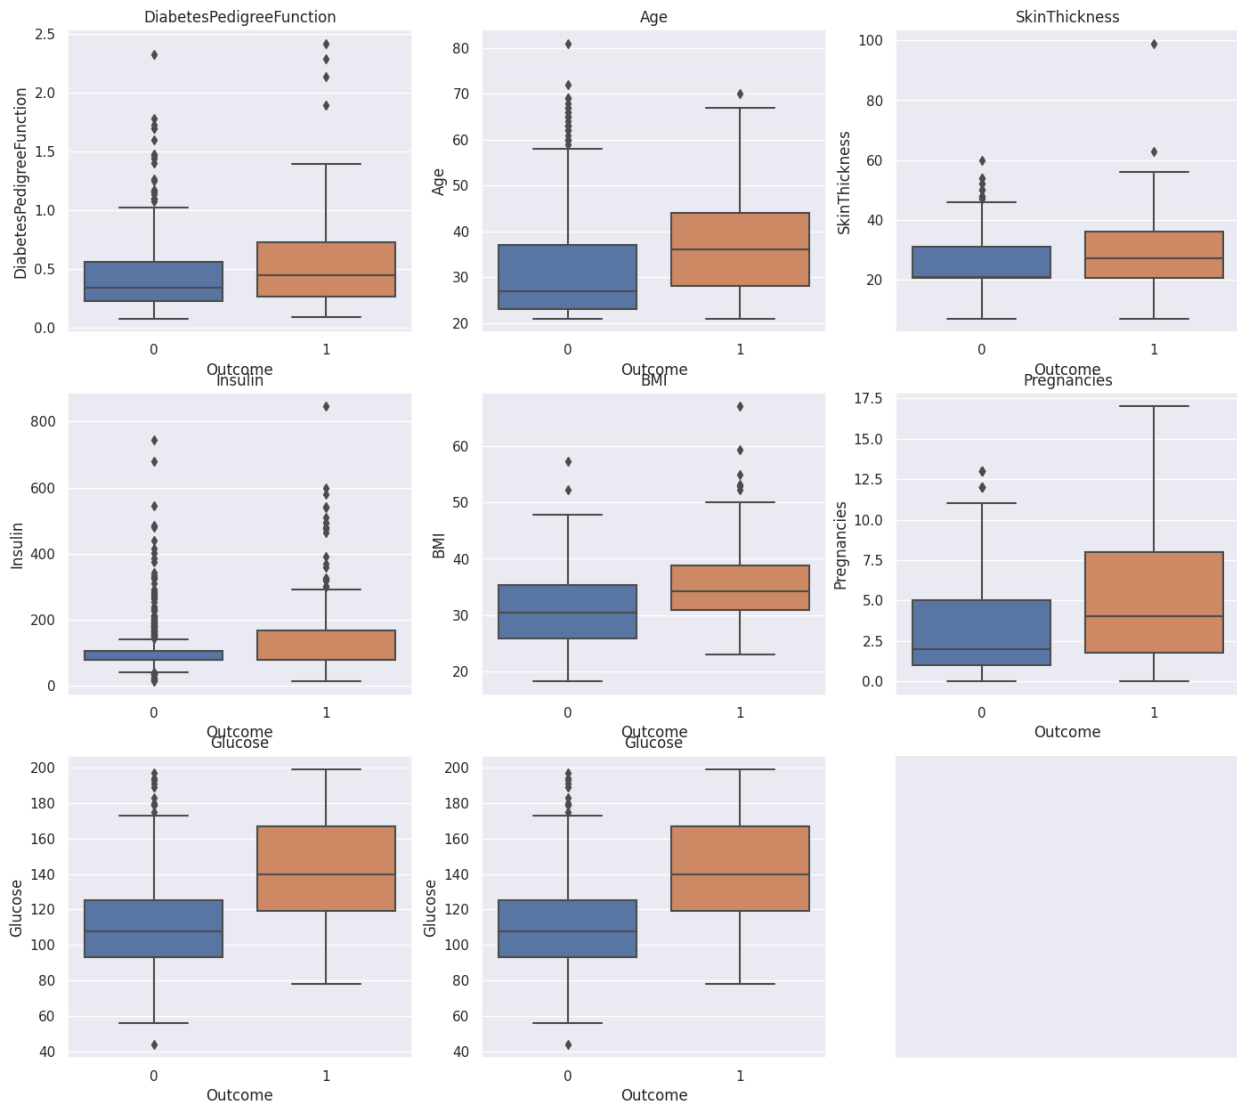


```

fi, ax = plt.subplots(3, 3)
fi.set_figheight(15)
fi.set_figwidth(17)
fi.add_subplot(3, 3, 1)
sns.boxplot(x = diabetes["Outcome"], y =
diabetes["DiabetesPedigreeFunction"])
plt.title("DiabetesPedigreeFunction")
fi.add_subplot(3, 3, 2)
sns.boxplot(x = diabetes["Outcome"], y = diabetes["Age"])
plt.title("Age")
fi.add_subplot(3, 3, 3)
sns.boxplot(x = diabetes["Outcome"], y = diabetes["SkinThickness"])
plt.title("SkinThickness")
fi.add_subplot(3, 3, 4)

```

```
sns.boxplot(x = diabetes["Outcome"], y = diabetes["Insulin"])
plt.title("Insulin")
fi.add_subplot(3, 3, 5)
sns.boxplot(x = diabetes["Outcome"], y = diabetes["BMI"])
plt.title("BMI")
fi.add_subplot(3, 3, 6)
sns.boxplot(x = diabetes["Outcome"], y = diabetes["Pregnancies"])
plt.title("Pregnancies")
fi.add_subplot(3, 3, 7)
sns.boxplot(x = diabetes["Outcome"], y = diabetes["Glucose"])
plt.title("Glucose")
fi.add_subplot(3, 3, 8)
sns.boxplot(x = diabetes["Outcome"], y = diabetes["Glucose"])
plt.title("Glucose")
for a in ax.flatten():
    a.set_xticks([])
    a.set_yticks([])
    a.set_xticklabels([])
    a.set_yticklabels([])
    a.spines['top'].set_visible(False)
    a.spines['right'].set_visible(False)
    a.spines['bottom'].set_visible(False)
    a.spines['left'].set_visible(False)
plt.show()
```



```

from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import precision_score, recall_score, f1_score
X = diabetes.iloc[:, :-1]
y = diabetes.iloc[:, -1:]
X_train, X_test, y_train, y_test1 = train_test_split(X, y,
test_size=0.3)
clf = DecisionTreeClassifier()
model = clf.fit(X_train, y_train)
y_pred1 = model.predict(X_test)
print("Tocnost (decision tree): ", metrics.accuracy_score(y_test1,
y_pred1)) #Njihova tocnost je oko 76%
print("Precision (decision tree): ", metrics.precision_score(y_test1,
y_pred1)) #Njima 70%

```



```

print(metrics.confusion_matrix(y_test1, y_pred1))
tn, fp, fn, tp = metrics.confusion_matrix(y_test1, y_pred1).ravel()
Spec1 = tn / (tn+fp)
Acc1 = metrics.accuracy_score(y_test1, y_pred1)
Prec1 = metrics.precision_score(y_test1, y_pred1)
Rec1 = metrics.recall_score(y_test1, y_pred1)
F1_1 = metrics.f1_score(y_test1, y_pred1)
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
print()
clf2 = DecisionTreeClassifier()
model2 = clf2.fit(X_train, y_train)
y_pred4 = model2.predict(X_test)
print("Tocnost (decision tree) (skalirano): ",
metrics.accuracy_score(y_test1, y_pred4))
print(metrics.confusion_matrix(y_test1, y_pred4))
tn, fp, fn, tp = metrics.confusion_matrix(y_test1, y_pred4).ravel()
Spec4 = tn / (tn+fp)
Acc4 = metrics.accuracy_score(y_test1, y_pred4)
Prec4 = metrics.precision_score(y_test1, y_pred4)
Rec4 = metrics.recall_score(y_test1, y_pred4)
F1_4 = metrics.f1_score(y_test1, y_pred4)

```

```

Tocnost (decision tree):  0.6926406926406926
Precision (decision tree):  0.5925925925925926
[[112  33]
 [ 38  48]]

```

```

Tocnost (decision tree) (skalirano):  0.658008658008658
[[106  39]
 [ 40  46]]

```

```

from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier()
X = diabetes.iloc[:, :-1]
y = diabetes.iloc[:, -1:]
X_train, X_test, y_train, y_test2 = train_test_split(X, y,
test_size=0.3)
model = clf.fit(X_train, y_train)
y_pred2 = model.predict(X_test)
print("Tocnost (random forest): ", metrics.accuracy_score(y_test2,
y_pred2)) #Njihova tocnost je oko 81%
print("Precision (random forest): ", metrics.precision_score(y_test2,
y_pred2)) #Njima 89%
print(metrics.confusion_matrix(y_test2, y_pred2))
tn, fp, fn, tp = metrics.confusion_matrix(y_test2, y_pred2).ravel()
Spec2 = tn / (tn+fp)
Acc2 = metrics.accuracy_score(y_test2, y_pred2)
Prec2 = metrics.precision_score(y_test2, y_pred2)

```

```

Rec2 = metrics.recall_score(y_test2, y_pred2)
F1_2 = metrics.f1_score(y_test2, y_pred2)
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
print()
clf2 = RandomForestClassifier()
model2 = clf2.fit(X_train, y_train)
y_pred5 = model2.predict(X_test)
print("Tocnost (random forest) (skalirano): ",
      metrics.accuracy_score(y_test2, y_pred5))
print(metrics.confusion_matrix(y_test2, y_pred5))
tn, fp, fn, tp = metrics.confusion_matrix(y_test2, y_pred5).ravel()
Spec5 = tn / (tn+fp)
Acc5 = metrics.accuracy_score(y_test2, y_pred5)
Prec5 = metrics.precision_score(y_test2, y_pred5)
Rec5 = metrics.recall_score(y_test2, y_pred5)
F1_5 = metrics.f1_score(y_test2, y_pred5)

```

```

<ipython-input-97-1e4957305201>:6: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the
shape of y to (n_samples,), for example using ravel().
    model = clf.fit(X_train, y_train)

```

```

Tocnost (random forest):  0.70995670995671
Precision (random forest):  0.6521739130434783
[[119  24]
 [ 43  45]]

```

```

<ipython-input-97-1e4957305201>:22: DataConversionWarning: A column-
vector y was passed when a 1d array was expected. Please change the
shape of y to (n_samples,), for example using ravel().
    model2 = clf2.fit(X_train, y_train)

```

```

Tocnost (random forest) (skalirano):  0.70995670995671
[[118  25]
 [ 42  46]]

```

```

from sklearn.naive_bayes import GaussianNB
X = diabetes.iloc[:, :-1]
y = diabetes.iloc[:, -1:]
X_train, X_test, y_train, y_test3 = train_test_split(X, y,
    test_size=0.3)
clf = GaussianNB()
model = clf.fit(X_train, y_train)
y_pred3 = model.predict(X_test)
print("Tocnost (naivni Bayes): ", metrics.accuracy_score(y_test3,
    y_pred3)) #Njihova tocnost je oko 78%
print("Precision (naivni Bayes): ", metrics.precision_score(y_test3,

```



```

y_pred3)) #Njima 81%
print(metrics.confusion_matrix(y_test3, y_pred3))
tn, fp, fn, tp = metrics.confusion_matrix(y_test3, y_pred3).ravel()
Spec3 = tn / (tn+fp)
Acc3 = metrics.accuracy_score(y_test3, y_pred3)
Prec3 = metrics.precision_score(y_test3, y_pred3)
Rec3 = metrics.recall_score(y_test3, y_pred3)
F1_3 = metrics.f1_score(y_test3, y_pred3)
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
print()
clf2 = GaussianNB()
model2 = clf2.fit(X_train, y_train)
y_pred6 = model2.predict(X_test)
print("Tocnost (naivni Bayes) (skalirano): ",
metrics.accuracy_score(y_test3, y_pred6))
print(metrics.confusion_matrix(y_test3, y_pred6))
tn, fp, fn, tp = metrics.confusion_matrix(y_test3, y_pred6).ravel()
Spec6 = tn / (tn+fp)
Acc6 = metrics.accuracy_score(y_test3, y_pred6)
Prec6 = metrics.precision_score(y_test3, y_pred6)
Rec6 = metrics.recall_score(y_test3, y_pred6)
F1_6 = metrics.f1_score(y_test3, y_pred6)

```

```

Tocnost (naivni Bayes): 0.7489177489177489
Precision (naivni Bayes): 0.704225352112676
[[123  21]
 [ 37  50]]

```

```

Tocnost (naivni Bayes) (skalirano): 0.7489177489177489
[[123  21]
 [ 37  50]]

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/
validation.py:1143: DataConversionWarning: A column-vector y was
passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:11
43: DataConversionWarning: A column-vector y was passed when a 1d
array was expected. Please change the shape of y to (n_samples, ), for
example using ravel().
    y = column_or_1d(y, warn=True)

```

```


from sklearn.metrics import roc_auc_score, roc_curve

fpr1, tpr1, thresholds1 = roc_curve(y_test1, y_pred1)
plt.plot(fpr1, tpr1, label="Decision Tree")
fpr2, tpr2, thresholds2 = roc_curve(y_test2, y_pred2)

```


```
plt.plot(fpr2, tpr2, label="Random Forest", color = "green")
fpr3, tpr3, thresholds3 = roc_curve(y_test3, y_pred3)
plt.plot(fpr3, tpr3, label="Naïve Bayes", color = "red")
plt.legend()
```

<matplotlib.legend.Legend at 0x7c959e847f10>



```
fpr4, tpr4, thresholds4 = roc_curve(y_test1, y_pred4)
plt.plot(fpr4, tpr4, label="Decision Tree (scaled)")
fpr5, tpr5, thresholds5 = roc_curve(y_test2, y_pred5)
plt.plot(fpr5, tpr5, label="Random Forest (scaled)", color = "green")
fpr6, tpr6, thresholds6 = roc_curve(y_test3, y_pred6)
plt.plot(fpr6, tpr6, label="Naïve Bayes (scaled)", color = "red")
plt.legend()
```

<matplotlib.legend.Legend at 0x7c959e8a5480>



```

fi, ax = plt.subplots(5, 2)
fi.set_figheight(25)
fi.set_figwidth(14)
fi.add_subplot(5, 2, 1)
plt.title("Accuracy")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Acc1,
Acc2, Acc3])
fi.add_subplot(5, 2, 2)
plt.title("Accuracy (scaled)")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Acc4,
Acc5, Acc6])
fi.add_subplot(5, 2, 3)
plt.title("Precision")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Prec1,
Prec2, Prec3])
fi.add_subplot(5, 2, 4)
plt.title("Precision (scaled)")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Prec4,
Prec5, Prec6])
fi.add_subplot(5, 2, 5)
plt.title("Recall")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Rec1,
Rec2, Rec3])
fi.add_subplot(5, 2, 6)
plt.title("Recall (scaled)")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Rec4,
Rec5, Rec6])

```

```

fi.add_subplot(5, 2, 7)
plt.title("Specificity")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Spec1,
Spec2, Spec3])
fi.add_subplot(5, 2, 8)
plt.title("Specificity (scaled)")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [Spec4,
Spec5, Spec6])
fi.add_subplot(5, 2, 9)
plt.title("F1 score")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [F1_1,
F1_2, F1_3])
fi.add_subplot(5, 2, 10)
plt.title("F1 score (scaled)")
plt.bar(["Decision Tree", "Random Forest", "Naive Bayes"], [F1_4,
F1_5, F1_6])

for a in ax.flatten():
    a.set_xticks([])
    a.set_yticks([])
    a.set_xticklabels([])
    a.set_yticklabels([])
    a.spines['top'].set_visible(False)
    a.spines['right'].set_visible(False)
    a.spines['bottom'].set_visible(False)
    a.spines['left'].set_visible(False)
plt.show()

```

